

Himanshu Pandey

✉ 26anjay06@gmail.com

☎ +91-7652023417

🌐 [linkedin.com/in/himanshu-pandey12](https://www.linkedin.com/in/himanshu-pandey12)

📄 github.com/Anzay12

EDUCATION

Indian Institute of Technology (IIT) Madras, Chennai, India

BS in Data Science and Applications

Sept. 2022 – Aug. 2026

Diploma in Programming & App Dev — Coursework: Data Structures, DBMS, Algorithms, Statistics

WORK EXPERIENCE

Shodhshala, Remote

Backend & Data Engineer Intern

Jan. 2026 – Present.

- Designed and maintained **PostgreSQL** databases (1M+ rows of high-frequency financial data) with query optimization, reducing average response time by ~30%.
- Built automated **ETL pipelines** in Python for large-scale trading datasets, cutting preprocessing time by ~40%.
- Engineered **30+ features** from raw market data via modular Python scripts, improving model signal quality by ~15%.
- Developed and backtested algorithmic trading strategies using statistical analysis pipelines, improving average win rate by ~12%.
- Monitored deployed system performance and delivered data-driven insights to refine live production strategies.

TECHNICAL PROJECTS

E-commerce RTO Prediction

Guide: github.com/Anzay12/Prediction_RTO Tech: Python, Gradient Boosting, Scikit-learn, Pandas

- Built a **Gradient Boosting classifier** achieving **68.75% F1-Score** and **73.33% ROC-AUC** to predict Return-to-Origin orders; benchmarked 5 models via cross-validation (CV: 70.66% ± 1.42%).
- Engineered 4 domain-specific features (Cost.Weight.Ratio, Weight.Bin, etc.) contributing **54%+ of total feature importance**; Weight and Discount identified as top RTO drivers.

Movie Audience Prediction — ML & Time Series

Guide: github.com/Anzay12/movie-audience-prediction-ml Tech: Python, XGBoost, LightGBM, Scikit-learn, Pandas

- Built an end-to-end forecasting pipeline with lag features (1/7/14/30-day) and rolling means across Random Forest, XGBoost, and LightGBM; selected best model via RMSE comparison.
- Engineered a weighted ensemble blending day-of-week patterns, 7 & 30-day trends, and theater-level behavior with multi-level fallback, improving prediction stability over single-model baseline.

Vehicle Parking System — Distributed Web Application

Guide: github.com/Anzay12 Tech: Python, Flask, SQLAlchemy, Celery, Redis, Vue.js, SQLite, REST APIs

- Architected a **role-based REST API backend** (Admin & User) using **Flask + SQLAlchemy** with real-time slot booking, reservation history, and authentication middleware.
- Implemented **Celery + Redis** for async background tasks — automated daily reminders and monthly CSV report generation via scheduled jobs, fully decoupled from the request cycle.
- Designed SQLite schema with **soft-deletion logic** to preserve historical reservation records, ensuring data integrity without hard deletes.

SKILLS

Backend: Python, Flask, FastAPI, REST APIs, SQLAlchemy, Celery, Redis, JWT Authentication

Databases: PostgreSQL, SQLite, MongoDB, Query Optimization, Schema Design, ETL Pipelines

DevOps: Docker, Kubernetes, CI/CD, Linux, Git

MLOps: MLflow, DVC, DagsHub, Grafana, Prometheus, Model Registry, Experiment Tracking, Data Versioning, Pipeline Automation

ML & Data: Scikit-learn, XGBoost, PyTorch, TensorFlow, Pandas, NumPy, EDA, Data Visualization

Frontend: React, Vue.js, JavaScript, HTML/CSS

CERTIFICATIONS

- Claude Code 101**
- MLOps Bootcamp: Mastering AI Operations**
- AWS Cloud Practitioner Essentials**
- Generative AI**

Anthropic
Udemy
GeeksforGeeks
Udemy